

Securing Docker on the Cheap

Part 1 - Fundamentals



About Me

- Possessor of many hats
- Currently at LO3 Energy
- Formerly of Autodesk
- This talk brought to you by the letter 'A'




But first....



The whale in the room



Orchestration platform agnostic



No experience
required.



Docker 101




What is a container?



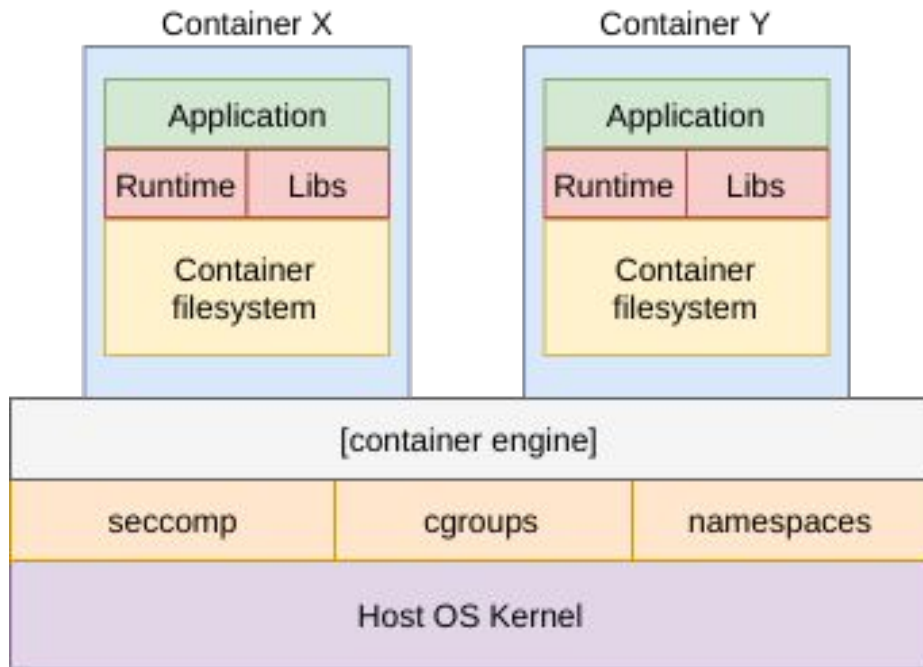
Perfect Packaging

- Everything needed for execution of the container is packaged together
 - Application code
 - Runtime environment
 - Dependency libraries
 - Even root filesystem
- Container “package” can be delivered to any host capable of executing that container



Intrinsic Isolation

- Pseudo-virtualization
- Isolated resources on top of a shared host kernel
- Container engine (LXC, Docker, etc.)
- Leverages control groups for resource allocation
- Namespaces provide isolation
- Has roots in chroot and FreeBSD jails



Why Docker?



Containers made easy



Adoption



Security Fundamentals



First...



Service Containers VS Tool Containers



Security Starts at the Top



Sample Dockerfile

```
FROM ubuntu:16.04

RUN apt update && apt upgrade -y && apt install -y curl && \
    curl -sL https://deb.nodesource.com/setup_8.x | bash - && \
    apt install -y nodejs

EXPOSE 3000

ADD app.js /var/app/
ADD package.json /var/app/

WORKDIR /var/app
RUN npm install


CMD ["/usr/bin/node", "app.js"]
```



Know your FROM



Security scan of ubuntu:16.04



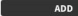







Quay Security Scanner has detected **44** vulnerabilities.

Patches are available for **5** vulnerabilities.

- 1 High-level vulnerabilities.
- 10 Medium-level vulnerabilities.
- 18 Low-level vulnerabilities.
- 15 Negligible-level vulnerabilities.


Image Vulnerabilities

☐ Only show fixable

| CVE | SEVERITY ↓ | PACKAGE | CURRENT VERSION | FIXED IN VERSION | INTRODUCED IN IMAGE |
|------------------------------------|--|----------|--------------------|---------------------|--|
| ▶ CVE-2018-10000 🔗 | 🚨 High | glibc | 2.23-0ubuntu10 | (None) |  file: 4c266e490f4101f9726598... |
| ▶ CVE-2017-8804 🔗 | 7.8 / 10  | glibc | 2.23-0ubuntu10 | (None) |  file: 4c266e490f4101f9726598... |
| ▶ CVE-2016-1238 🔗 | 7.2 / 10  | perl | 5.22.1-9ubuntu0.2 | (None) |  file: 4c266e490f4101f9726598... |
| ▶ CVE-2018-6485 🔗 | 🚨 Medium | glibc | 2.23-0ubuntu10 | (None) |  file: 4c266e490f4101f9726598... |
| ▶ CVE-2016-1585 🔗 | 🚨 Medium | apparmor | 2.10.95-0ubuntu2.9 | (None) |  file: 4c266e490f4101f9726598... |
| ▶ CVE-2018-6913 🔗 | 🚨 Medium | perl | 5.22.1-9ubuntu0.2 | 🟢 5.22.1-9ubuntu0.3 |  file: 4c266e490f4101f9726598... |



Security scan of node:9.11.1



Quay Security Scanner has detected **634** vulnerabilities.

Patches are available for **6** vulnerabilities.

- 83** High-level vulnerabilities.
- 212** Medium-level vulnerabilities.
- 119** Low-level vulnerabilities.
- 192** Negligible-level vulnerabilities.
- 28** Unknown-level vulnerabilities.

Image Vulnerabilities

Filter Vulnerabilities... ☐ Only show fixable

| CVE | SEVERITY ↓ | PACKAGE | CURRENT VERSION | FIXED IN VERSION | INTRODUCED IN IMAGE |
|------------------|---------------------------------|-----------|----------------------|------------------|--------------------------------|
| ▶ CVE-2017-17458 | 10 / 10 <div><div></div></div> | mercurial | 3.1.2-2+deb8u4 | (None) | apt-get update && apt-get i... |
| ▶ CVE-2017-18017 | 10 / 10 <div><div></div></div> | linux | 3.16.51-3+deb8u1 | (None) | set -ex; apt-get update; ap... |
| ▶ CVE-2016-4448 | 10 / 10 <div><div></div></div> | libxml2 | 2.9.1+dfsg1-5+deb8u6 | (None) | set -ex; apt-get update; ap... |
| ▶ CVE-2015-1418 | 9.3 / 10 <div><div></div></div> | patch | 2.7.5-1 | (None) | set -ex; apt-get update; ap... |
| ▶ CVE-2017-16997 | 9.3 / 10 <div><div></div></div> | glibc | 2.19-18+deb8u10 | (None) | file:bc844c4763367b5f0ac7b9... |
| ▶ CVE-2016-3857 | 9.3 / 10 <div><div></div></div> | linux | 3.16.51-3+deb8u1 | (None) | set -ex; apt-get update; ap... |



Creating a Custom Base Container



Rules for a Quality Custom Base

- Starting tiny is better
- Patch as part of the build
- Build a shared service base
 - Install common tools
 - Install base runtime
- Leave the application specifics for downstream containers
- Leverage any hardening standards/tools for the OS
- Install only what you need
- Leave build tools on build containers

Pro-tip: `docker run -it [base-image]:[tag] /bin/sh` to experiment



Scratch Containers

- Docker images can be derived from tarballs
- Docker containers need a filesystem
- ...but that does not need to be a full base OS filesystem
- Docker images can be built directly
- FROM scratch
- Statically-linked executables (like Go) can be built directly into Docker images



Capabilities



Removing Capabilities



```
docker run --cap-drop SETUID --cap-drop CHOWN ubuntu
```



Adding Capabilities



```
docker run --cap-add RAW_IO
```



<https://github.com/moby/moby/blob/master/oci/defaults.go>

```
func defaultCapabilities() []string {  
    return []string{  
        "CAP_CHOWN", "CAP_DAC_OVERRIDE", "CAP_FSETID",  
        "CAP_FOWNER", "CAP_MKNOD", "CAP_NET_RAW",  
        "CAP_SETGID", "CAP_SETUID", "CAP_SETFCAP",  
        "CAP_SETPCAP", "CAP_NET_BIND_SERVICE",  
        "CAP_SYS_CHROOT", "CAP_KILL", "CAP_AUDIT_WRITE",  
    }  
}
```



Secure Computing Mode (seccomp)



<https://github.com/moby/moby/blob/master/profiles/seccomp/default.json>

```
{
  "defaultAction": "SCMP_ACT_ERRNO",
  "archMap": [],
  "syscalls": [
    {
      "names": [ "chdir", "chmod", ... ],
      "action": "SCMP_ACT_ALLOW",
      "args": [],
      "comment": "",
      "includes": {},
      "excludes": {}
    }
  ]
}
```




```
docker run --security-opt  
seccomp=/path/to/profile.json ubuntu: 16.04
```



Immutable Containers

```
docker run --read-only \  
--tmpfs /tmp:rw,noexec,nosuid \  
-v /host/path:/app/workdir [image]
```



Control Groups (cgroups)



Ok, what are control groups?



Docker & cgroups



Some cgroup Options

- cpu*
- blkio*
- device*
- memory*

<https://docs.docker.com/engine/reference/commandline/run/#options>



Other Tips



Don't run as root



Avoid `--privileged`



Be intentional



Wrap-up

<https://github.com/fork4/lfnw2018>

